

CSCI 330

The UNIX System



User Datagram Protocol

Unit Overview

- Transport layer
- User datagram protocol
- UDP programming

Network Layer

- also called: Internet Protocol Layer
 - provides host to host transmission service,
where hosts are not necessarily adjacent
- layer provides services:
 - addressing
 - hosts have global addresses: IPv4, IPv6
 - routing and forwarding
 - find path from host to host

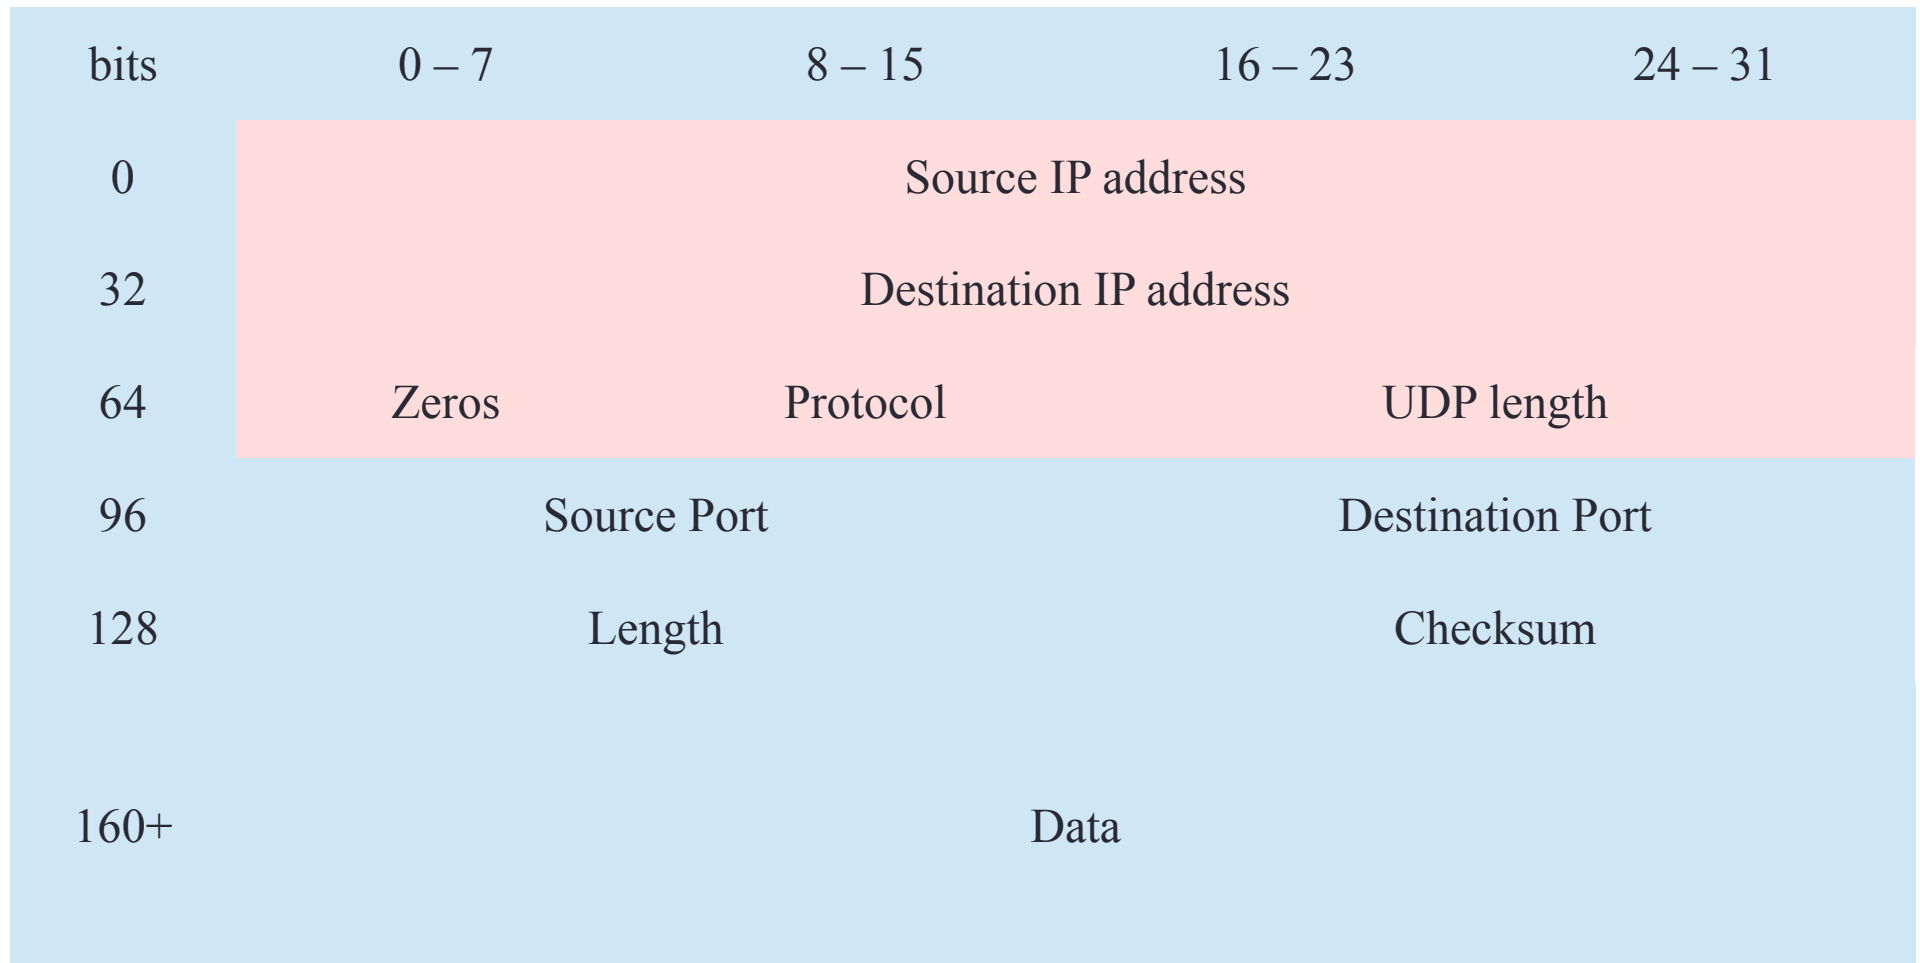
Transport Layer

- provides end-to-end communication services for applications
- provides multiple endpoints on a single node: port
- TCP: transmission control protocol
 - connection oriented, guaranteed delivery
 - stream oriented: basis for: http, ftp, smtp, ssh
- UDP: user datagram protocol
 - best effort
 - datagram oriented: basis for: dns, rtp

UDP

- simple message-based connectionless protocol
 - transmits information in one direction from source to destination without verifying the readiness or state of the receiver
- uses datagram as message
- stateless and fast

UDP packet format



UDP programming

- common abstraction: socket
- first introduced in BSD Unix in 1981
- socket is end-point of communication link
 - identified as IP address + port number
- operates as client and server

Socket system calls

server

System call

Meaning

client

socket

Create a new communication endpoint

bind

Attach a local address to a socket

sendto

Send(write) some data over the connection

recvfrom

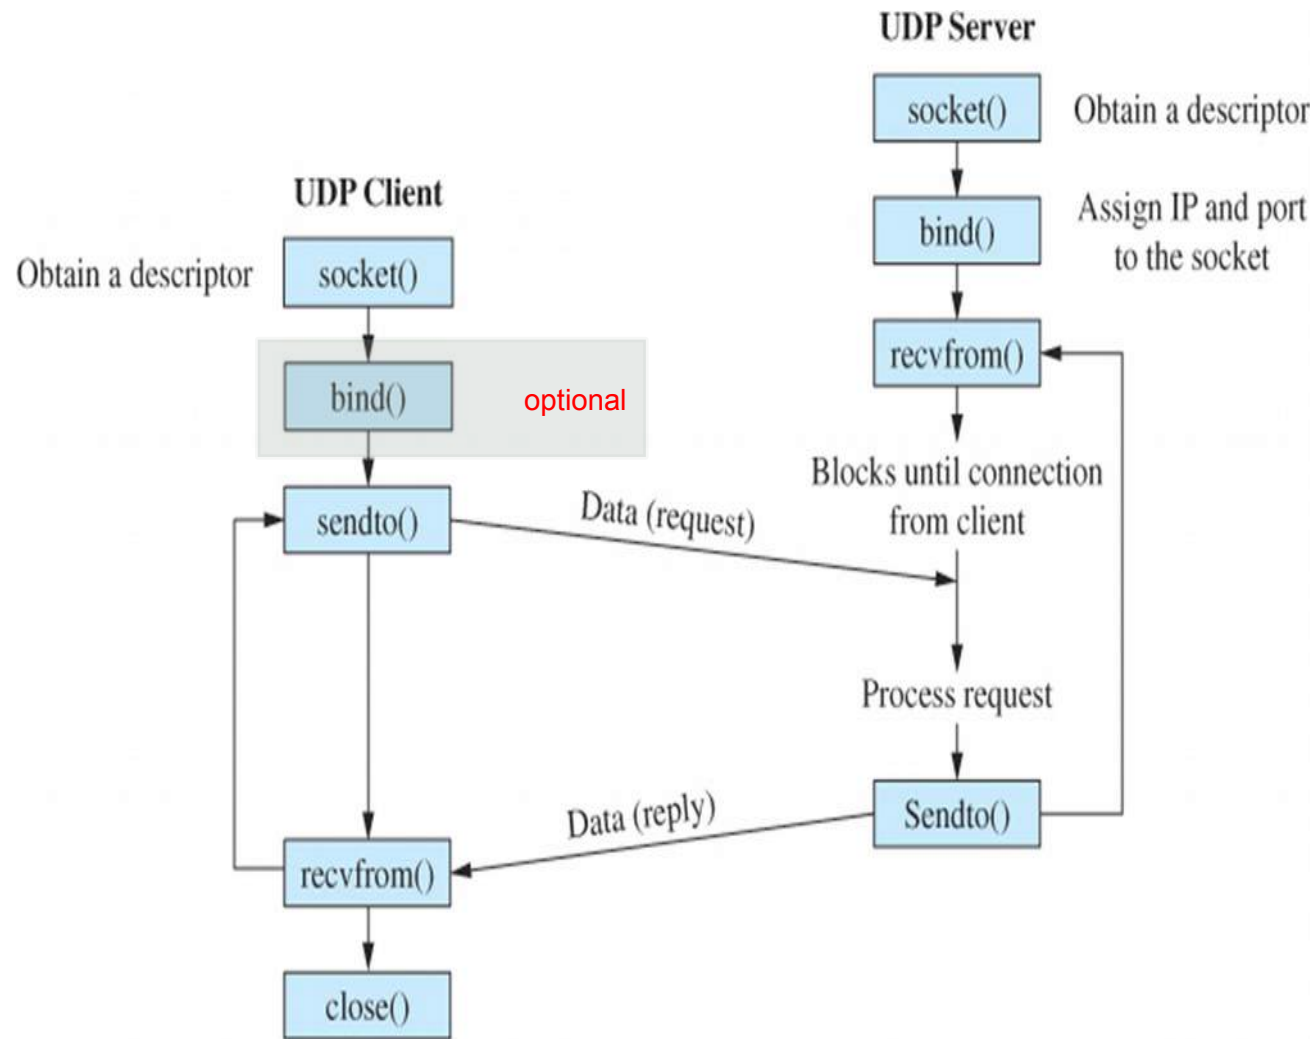
Receive(read) some data over the connection

close

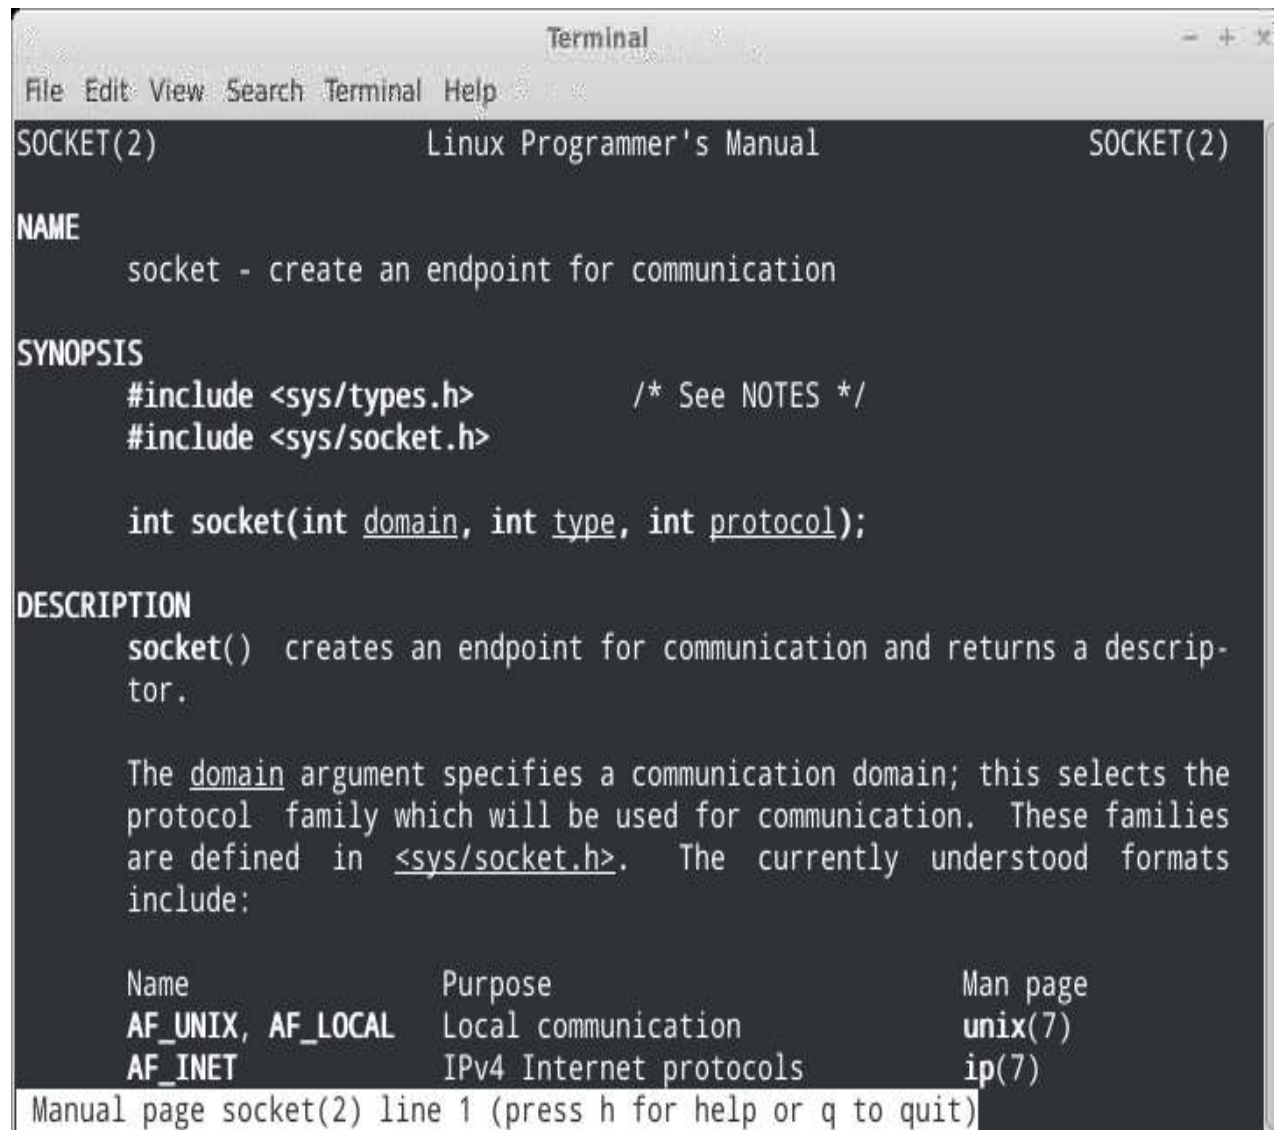
Release the connection

optional

UDP communications pattern



System call: socket



```
Terminal
File Edit View Search Terminal Help
SOCKET(2)          Linux Programmer's Manual          SOCKET(2)

NAME
    socket - create an endpoint for communication

SYNOPSIS
    #include <sys/types.h>          /* See NOTES */
    #include <sys/socket.h>

    int socket(int domain, int type, int protocol);

DESCRIPTION
    socket() creates an endpoint for communication and returns a descrip-
    tor.

    The domain argument specifies a communication domain; this selects the
    protocol family which will be used for communication. These families
    are defined in <sys/socket.h>. The currently understood formats
    include:

    Name                Purpose                Man page
    AF_UNIX, AF_LOCAL   Local communication    unix(7)
    AF_INET              IPv4 Internet protocols ip(7)

Manual page socket(2) line 1 (press h for help or q to quit)
```

System call: socket

```
int socket(int domain, int type, int protocol)
```

- creates a new socket, as end point to a communications link
- **domain** is set to **AF_INET**
- **type** is set to **SOCK_DGRAM** for datagrams
- **protocol** is set to 0, i.e. default UDP
- returns socket descriptor:
 - used in bind, sendto, recvfrom, close

System call: bind

```
Terminal
File Edit View Search Terminal Help
BIND(2) Linux Programmer's Manual BIND(2)

NAME
    bind - bind a name to a socket

SYNOPSIS
    #include <sys/types.h>          /* See NOTES */
    #include <sys/socket.h>

    int bind(int sockfd, const struct sockaddr *addr,
              socklen_t addrlen);

DESCRIPTION
    When a socket is created with socket(2), it exists in a name space
    (address family) but has no address assigned to it. bind() assigns
    the address specified to by addr to the socket referred to by the file
    descriptor sockfd. addrlen specifies the size, in bytes, of the
    address structure pointed to by addr. Traditionally, this operation
    is called "assigning a name to a socket".

    It is normally necessary to assign a local address using bind() before
    a SOCK_STREAM socket may receive connections (see accept(2)).

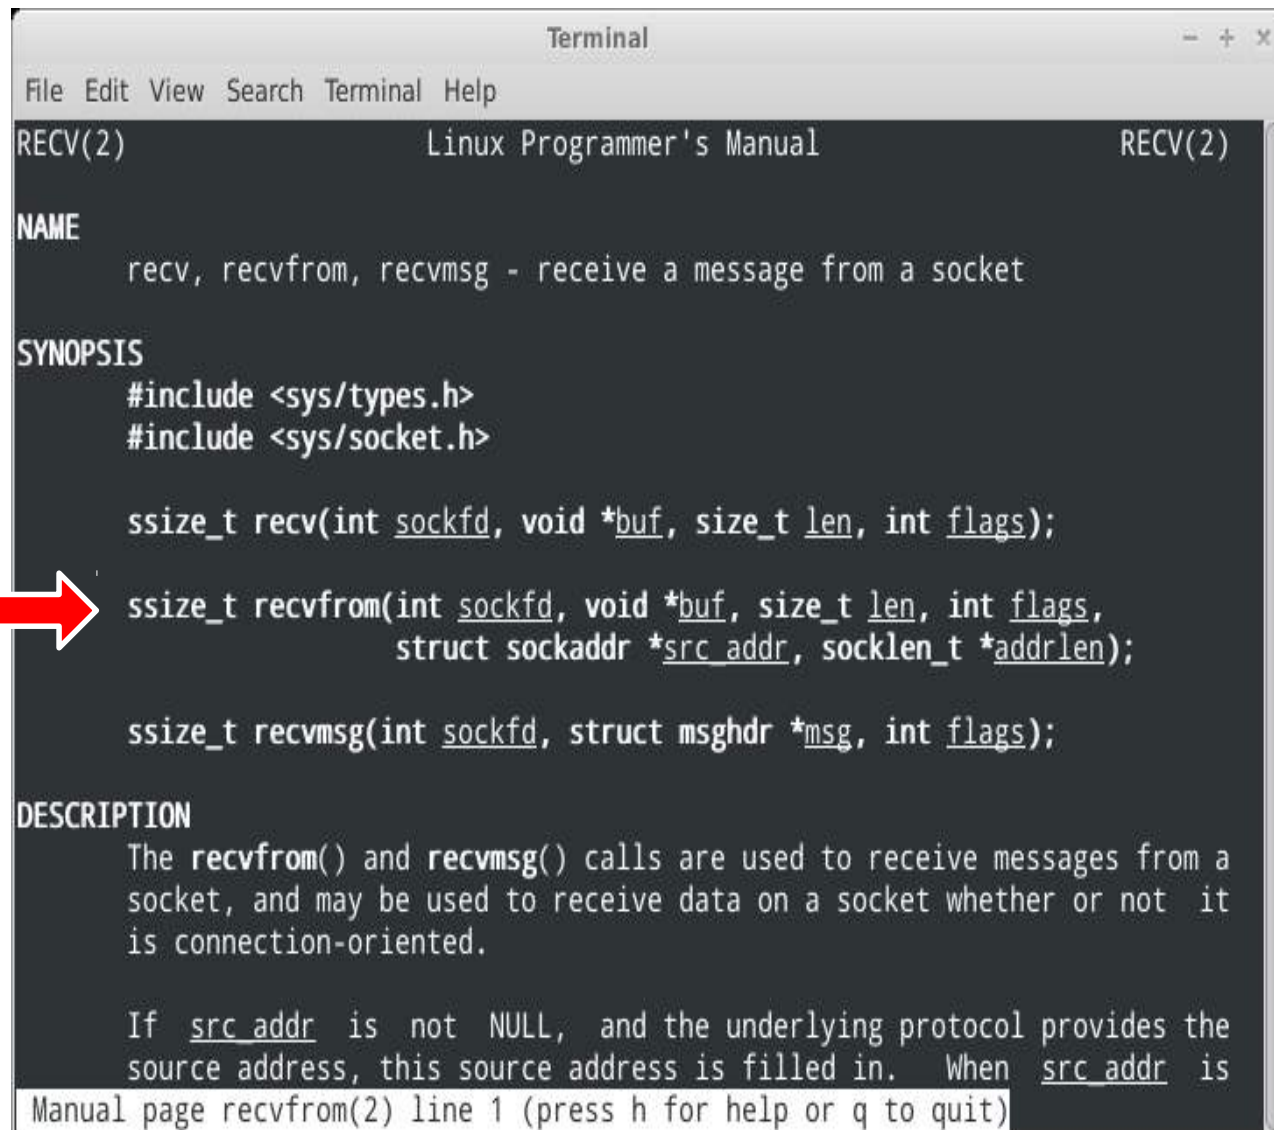
Manual page bind(2) line 1 (press h for help or q to quit)
```

System call: bind

```
int bind(int sockfd,  
         const struct sockaddr *addr,  
         socklen_t addrlen)
```

- assigns address to socket: IP number and port
- `struct sockaddr` holds address information
 - will accept `struct sockaddr_in`
- `addrlen` specifies length of `addr` structure
- returns 0 on success, -1 otherwise

System call: recvfrom



```
Terminal
File Edit View Search Terminal Help
recv(2) Linux Programmer's Manual recv(2)

NAME
recv, recvfrom, recvmsg - receive a message from a socket

SYNOPSIS
#include <sys/types.h>
#include <sys/socket.h>

ssize_t recv(int sockfd, void *buf, size_t len, int flags);
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags,
                 struct sockaddr *src_addr, socklen_t *addrlen);
ssize_t recvmsg(int sockfd, struct msghdr *msg, int flags);

DESCRIPTION
The recvfrom() and recvmsg() calls are used to receive messages from a
socket, and may be used to receive data on a socket whether or not it
is connection-oriented.

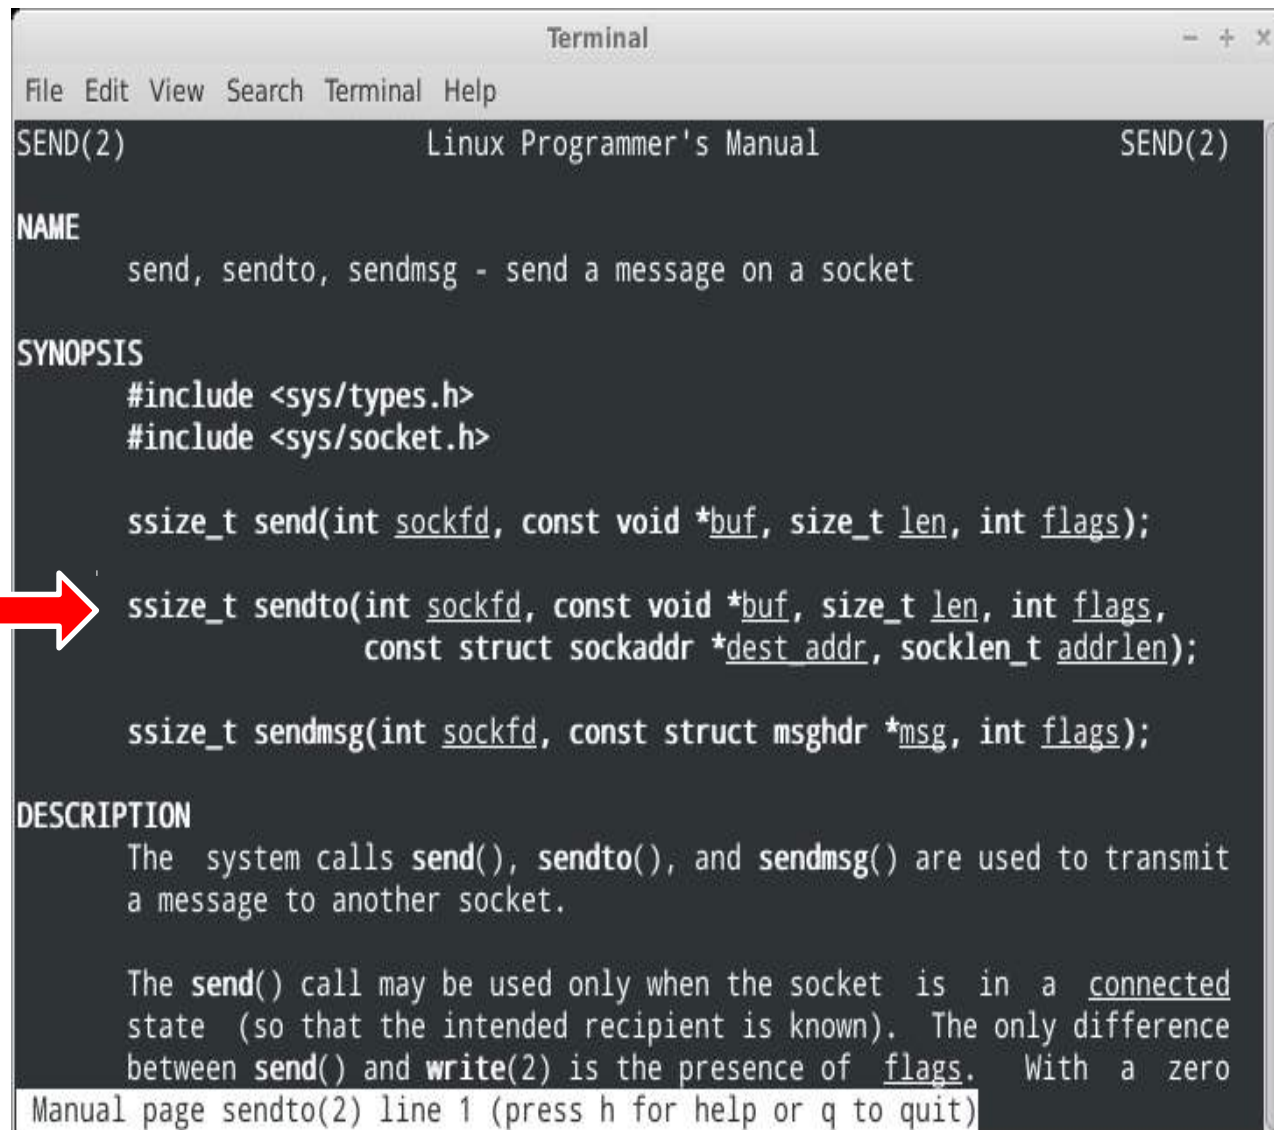
If src_addr is not NULL, and the underlying protocol provides the
source address, this source address is filled in. When src_addr is
Manual page recvfrom(2) line 1 (press h for help or q to quit)
```

System call: recvfrom

```
ssize_t recvfrom(int sockfd, void *buf, size_t len,  
                 int flags, struct sockaddr *src_addr,  
                 socklen_t *addrlen)
```

- receives a datagram **buf** of size **len** from socket
 - will wait until a datagram is available
 - flags specifies wait behavior, e.g.: 0 for default
- **src_addr** will hold address information of sender
 - **struct sockaddr** holds address information
 - **addrlen** specifies length of **src_addr** structure
- returns the number of bytes received, i.e. size of datagram

System call: sendto



```
Terminal
File Edit View Search Terminal Help
SEND(2) Linux Programmer's Manual SEND(2)

NAME
    send, sendto, sendmsg - send a message on a socket

SYNOPSIS
    #include <sys/types.h>
    #include <sys/socket.h>

    ssize_t send(int sockfd, const void *buf, size_t len, int flags);
    ssize_t sendto(int sockfd, const void *buf, size_t len, int flags,
                   const struct sockaddr *dest_addr, socklen_t addrlen);
    ssize_t sendmsg(int sockfd, const struct msghdr *msg, int flags);

DESCRIPTION
    The system calls send(), sendto(), and sendmsg() are used to transmit
    a message to another socket.

    The send() call may be used only when the socket is in a connected
    state (so that the intended recipient is known). The only difference
    between send() and write(2) is the presence of flags. With a zero
    Manual page sendto(2) line 1 (press h for help or q to quit)
```


System call: sendto

```
ssize_t sendto(int sockfd,  
               const void *buf, size_t len, int flags,  
               const struct sockaddr *dest_addr, socklen_t  
               addrlen)
```

- sends datagram **buf** of size **len** to socket
 - will wait if there is no ready receiver
 - flags specifies wait behavior, e.g.: 0 for default
- **dest_addr** holds address information of receiver
 - **struct sockaddr** holds address information
 - **addrlen** specifies length of **dest_addr** structure

UDP Programming

- simple server: echo
 - sends all received datagrams back to sender
- simple client
 - send datagram to server

Illustration: echoServer.cc

```
int sock;
struct sockaddr_in echoserver; // structure for address of server
struct sockaddr_in echoclient; // structure for address of client

// Create the UDP socket
if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
    perror("Failed to create socket"); exit(EXIT_FAILURE); }

// Construct the server sockaddr_in structure
memset(&echoserver, 0, sizeof(echoserver)); // Clear struct
echoserver.sin_family = AF_INET; // Internet IP
echoserver.sin_addr.s_addr = INADDR_ANY; // Any IP address
echoserver.sin_port = htons(atoi(argv[1])); // server port

// Bind the socket
serverlen = sizeof(echoserver);
if (bind(sock, (struct sockaddr *) &echoserver, serverlen) < 0) {
    perror("Failed to bind server socket"); exit(EXIT_FAILURE); }

// Run until cancelled
while (true) {
    // Receive a message from the client
    clientlen = sizeof(echoclient);
    if ((received = recvfrom(sock, buffer, 256, 0, (struct sockaddr *) &echoclient, &clientlen)) < 0) {
        perror("Failed to receive message"); exit(EXIT_FAILURE); }
    cerr << "Client connected: " << inet_ntoa(echoclient.sin_addr) << "\n";
    // Send the message back to client
    if (sendto(sock, buffer, received, 0, (struct sockaddr *) &echoclient, clientlen) != received) {
        perror("Mismatch in number of echo'd bytes"); exit(EXIT_FAILURE); }
}

close(sock);
```

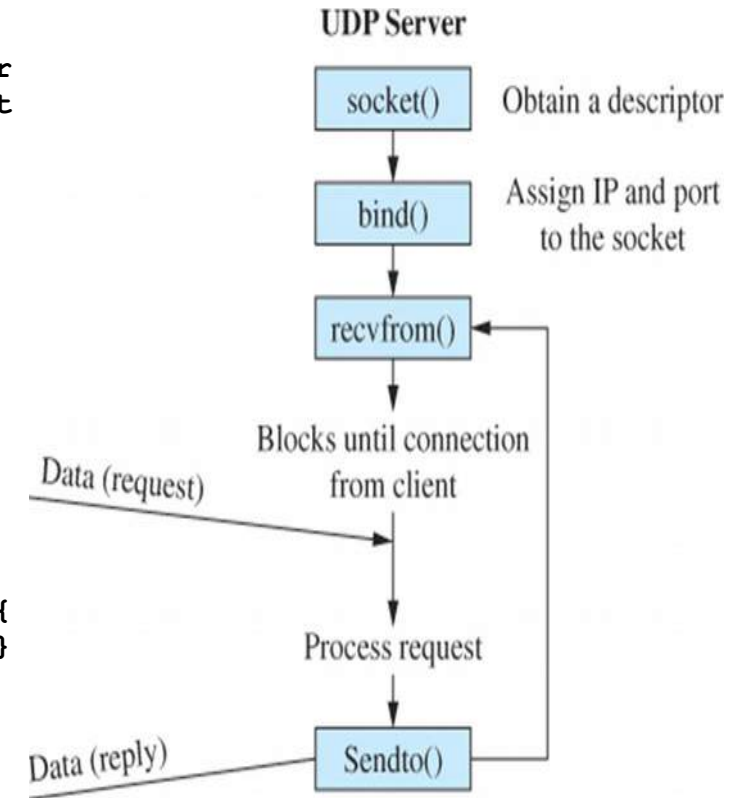


Illustration: echoClient.cc

```
int main(int argc, char * argv[])
{
    if(argc != 4) {
        cerr << "Usage: echoClient server_ip port message\n";
        exit(EXIT_FAILURE);
    }
    char buffer[256]; int echolen, received = 0; unsigned int addrlen;
    int sock;
    struct sockaddr_in echoserver; // structure for server address

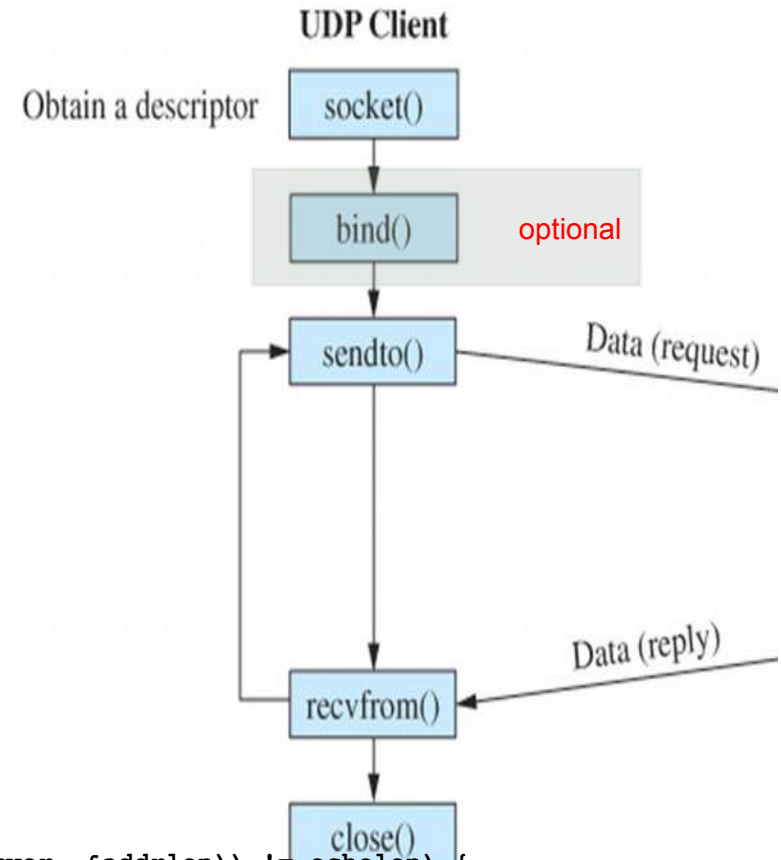
    // Create the UDP socket
    if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
        perror("Failed to create socket"); exit(EXIT_FAILURE); }

    // Construct the server sockaddr_in structure
    memset(&echoserver, 0, sizeof(echoserver)); // Clear struct
    echoserver.sin_family = AF_INET; // Internet IP
    echoserver.sin_addr.s_addr = inet_addr(argv[1]); // IP address
    echoserver.sin_port = htons(atoi(argv[2])); // server port

    // Send the message to the server
    echolen = strlen(argv[3]);
    if (sendto(sock, argv[3], strlen(argv[3]), 0,
        (struct sockaddr *) &echoserver, sizeof(echoserver)) != echolen) {
        perror("Mismatch in number of sent bytes"); exit(EXIT_FAILURE); }

    // Receive the message back from the server
    addrlen = sizeof(echoserver);
    if ((received = recvfrom(sock, buffer, 256, 0, (struct sockaddr *) &echoserver, &addrlen)) != echolen) {
        perror("Mismatch in number of received bytes"); exit(EXIT_FAILURE); }

    buffer[received] = '\0';
    cout << "Message received: " << buffer << endl;
}
```



Detail: create UDP socket

```
int sock;  
  
// Create the UDP socket  
  
if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {  
    perror("Failed to create socket");  
    exit(EXIT_FAILURE);  
}
```

Detail: bind the socket

```
struct sockaddr_in echoserver; // structure for  
address of server
```

```
// Construct the server sockaddr_in structure
```

```
memset(&echoserver, 0, sizeof(echoserver));          /*  
Clear struct */
```

```
echoserver.sin_family = AF_INET;                     /*  
Internet/IP */
```

```
echoserver.sin_addr.s_addr = INADDR_ANY;             /*  
Any IP address */
```

```
echoserver.sin_port = htons(atoi(argv[1]));        /*  
server port */
```

Detail: receive from socket

```
addrlen = sizeof(echoserver);

received = recvfrom(sock, buffer, 256, 0,
                    (struct sockaddr *) &echoserver, &addrlen);

cout << "Received: << received bytes\n";

buffer[received] = '\0';    /* Assure null-terminated string */

cout << "Server ("
      << inet_ntoa(echoserver.sin_addr)
      << ") echoed: " << buffer << endl;
```

Detail: send to socket

```
// Construct the server sockaddr_in structure

memset(&echoserver, 0, sizeof(echoserver));      /* Clear struct */

echoserver.sin_family = AF_INET;                 /* Internet/IP */

echoserver.sin_addr.s_addr = inet_addr(argv[1]); /* IP address */

echoserver.sin_port = htons(atoi(argv[2]));     /* server port */


// Send the message to the server

echolen = strlen(argv[3]);

if (sendto(sock, argv[3], strlen(argv[3]), 0,

    (struct sockaddr *) &echoserver, sizeof(echoserver))

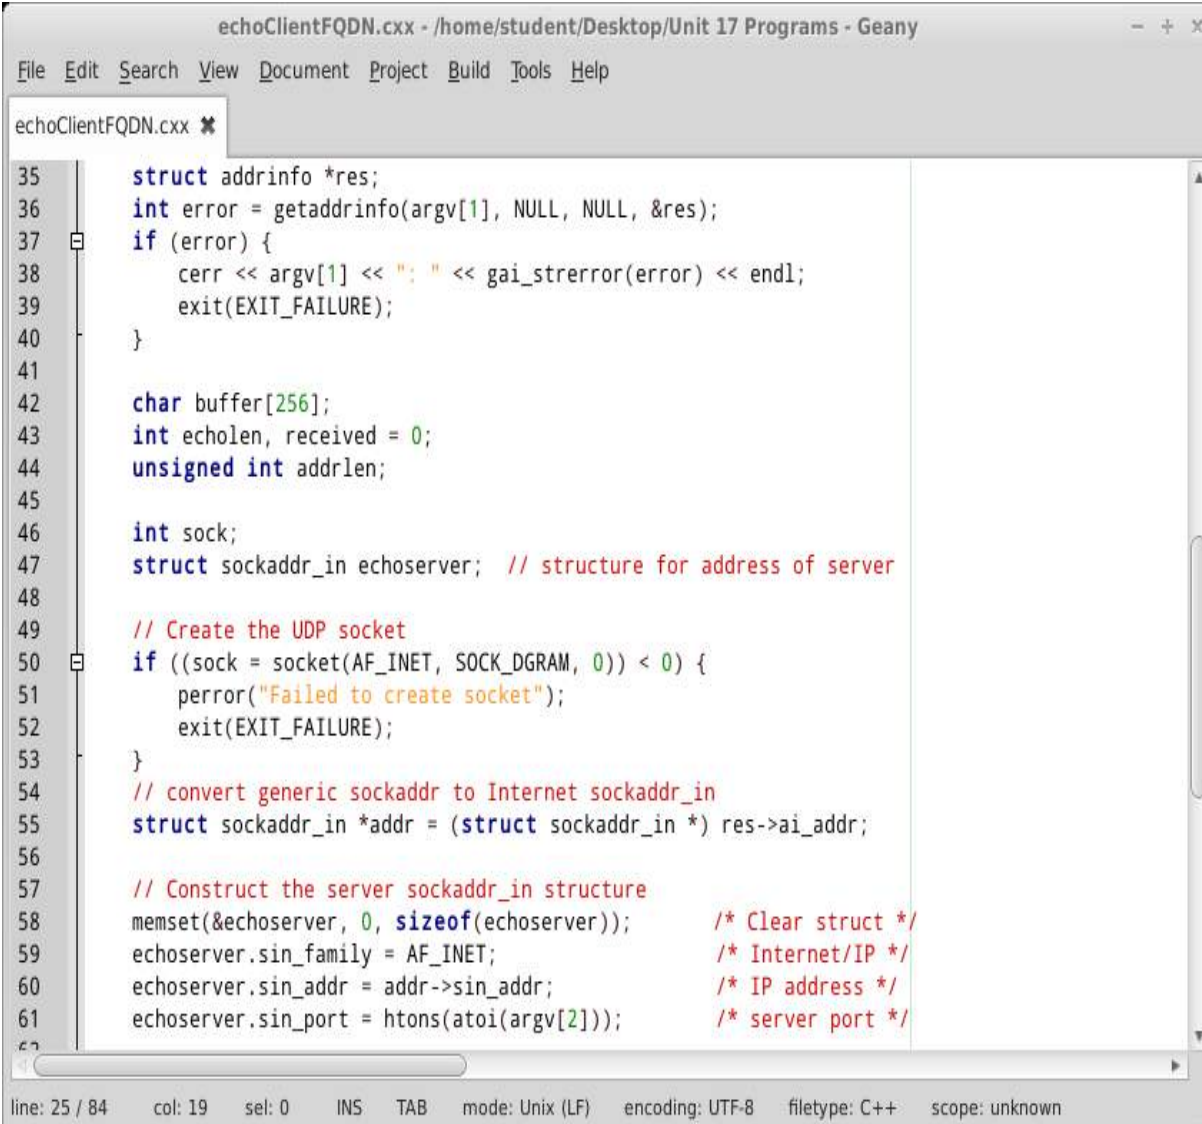
    != echolen) {

    perror("Mismatch in number of sent bytes");

    exit(EXIT_FAILURE);

}
```


Illustration: echoClientFQDN.cxx



```
echoClientFQDN.cxx - /home/student/Desktop/Unit 17 Programs - Geany
File Edit Search View Document Project Build Tools Help
echoClientFQDN.cxx *
35 struct addrinfo *res;
36 int error = getaddrinfo(argv[1], NULL, NULL, &res);
37 if (error) {
38     cerr << argv[1] << ": " << gai_strerror(error) << endl;
39     exit(EXIT_FAILURE);
40 }
41
42 char buffer[256];
43 int echolen, received = 0;
44 unsigned int addrlen;
45
46 int sock;
47 struct sockaddr_in echoserver; // structure for address of server
48
49 // Create the UDP socket
50 if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
51     perror("Failed to create socket");
52     exit(EXIT_FAILURE);
53 }
54 // convert generic sockaddr to Internet sockaddr_in
55 struct sockaddr_in *addr = (struct sockaddr_in *) res->ai_addr;
56
57 // Construct the server sockaddr_in structure
58 memset(&echoserver, 0, sizeof(echoserver)); /* Clear struct */
59 echoserver.sin_family = AF_INET; /* Internet/IP */
60 echoserver.sin_addr = addr->sin_addr; /* IP address */
61 echoserver.sin_port = htons(atoi(argv[2])); /* server port */
62
```

line: 25 / 84 col: 19 sel: 0 INS TAB mode: Unix (LF) encoding: UTF-8 filetype: C++ scope: unknown

Summary

- Transport layer
- User datagram protocol
- UDP programming