

CSCI 330

The UNIX System



Socket Programming Detail

Unit Overview

- TCP programming
- socket behavior
 - blocking vs. non-blocking
- signal handler

Socket system calls

server

Primitive

Meaning

client



socket

Create a new communication endpoint

bind

Attach a local address to a socket

listen

Announce willingness to accept connections

accept

Block caller until a connection request arrives

connect

Actively attempt to establish a connection

write

Send(write) some data over the connection

read

Receive(read) some data over the connection

close

Release the connection

TCP socket programming

- sockets can operate in one of two modes:
blocking or non-blocking
- default: blocking mode
 - calls may block, i.e. wait for something before they return
 - blocking is a special state in which a process is waiting for I/O
 - process is removed from the scheduler queue
 - when the I/O completes, the process is woken up
- non-blocking mode
 - calls returns immediately, producing full, partial or no result

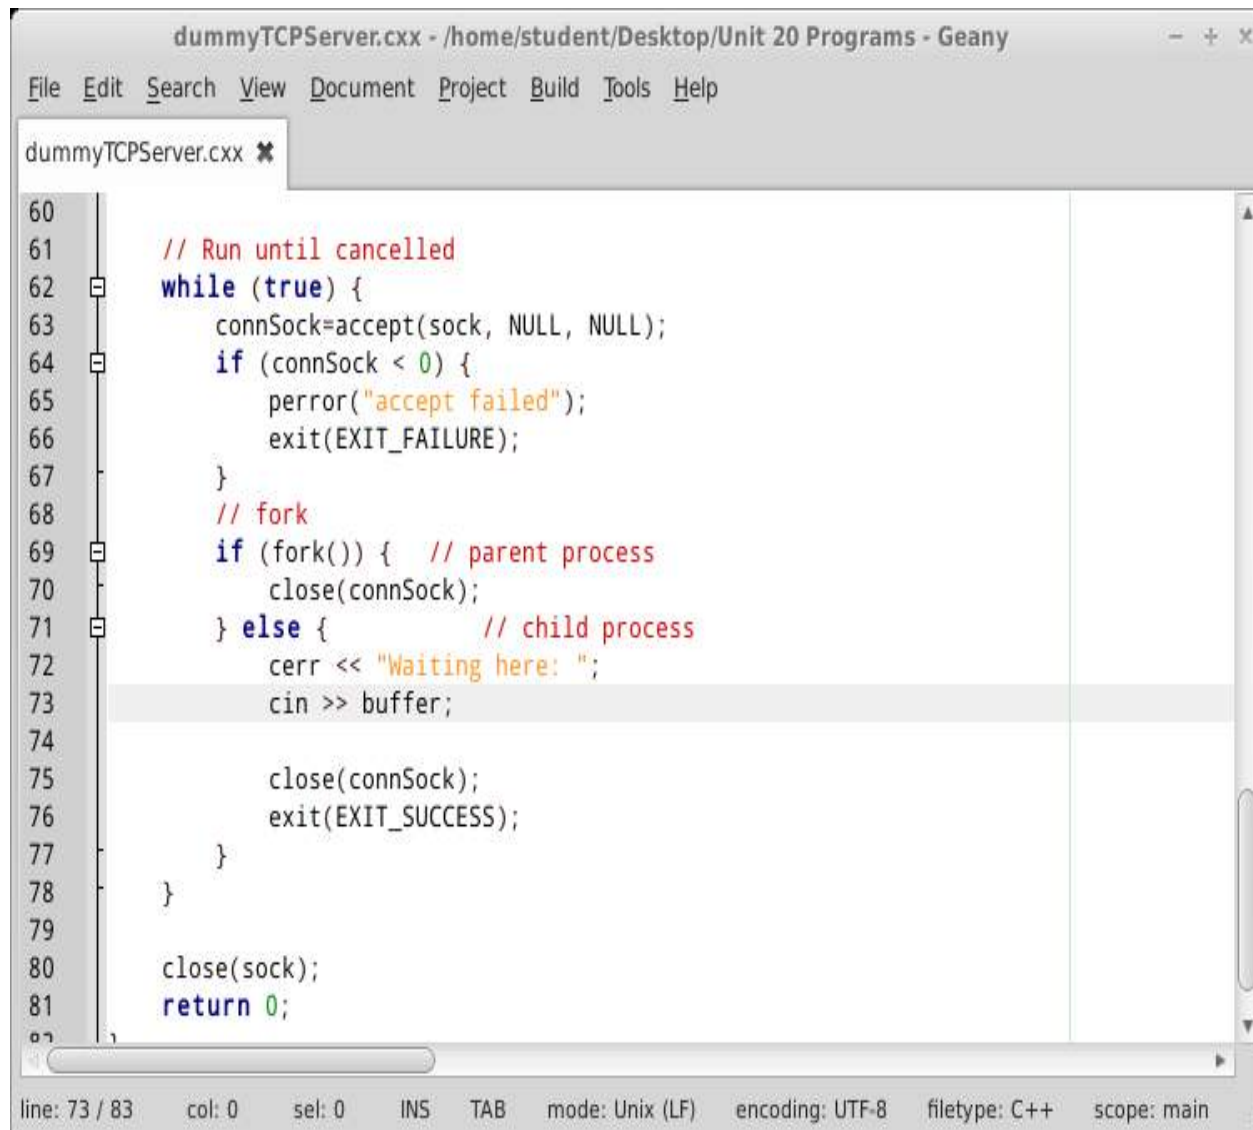
Blocking mode: calls that can block

- **accept**
 - blocks until a connection is present
- **connect**
 - blocks until connection is established
- **read**
 - blocks if no data is available
- **write**
 - blocks when data does not fit into send buffer of the socket

Illustration: blocking mode

```
// loop to write a single character
char c = 'x';
while (true) {
    if (write(sock, &c, 1) < 1) {
        perror("write");
        exit(EXIT_FAILURE);
    }
}
```

Illustration: dummy server

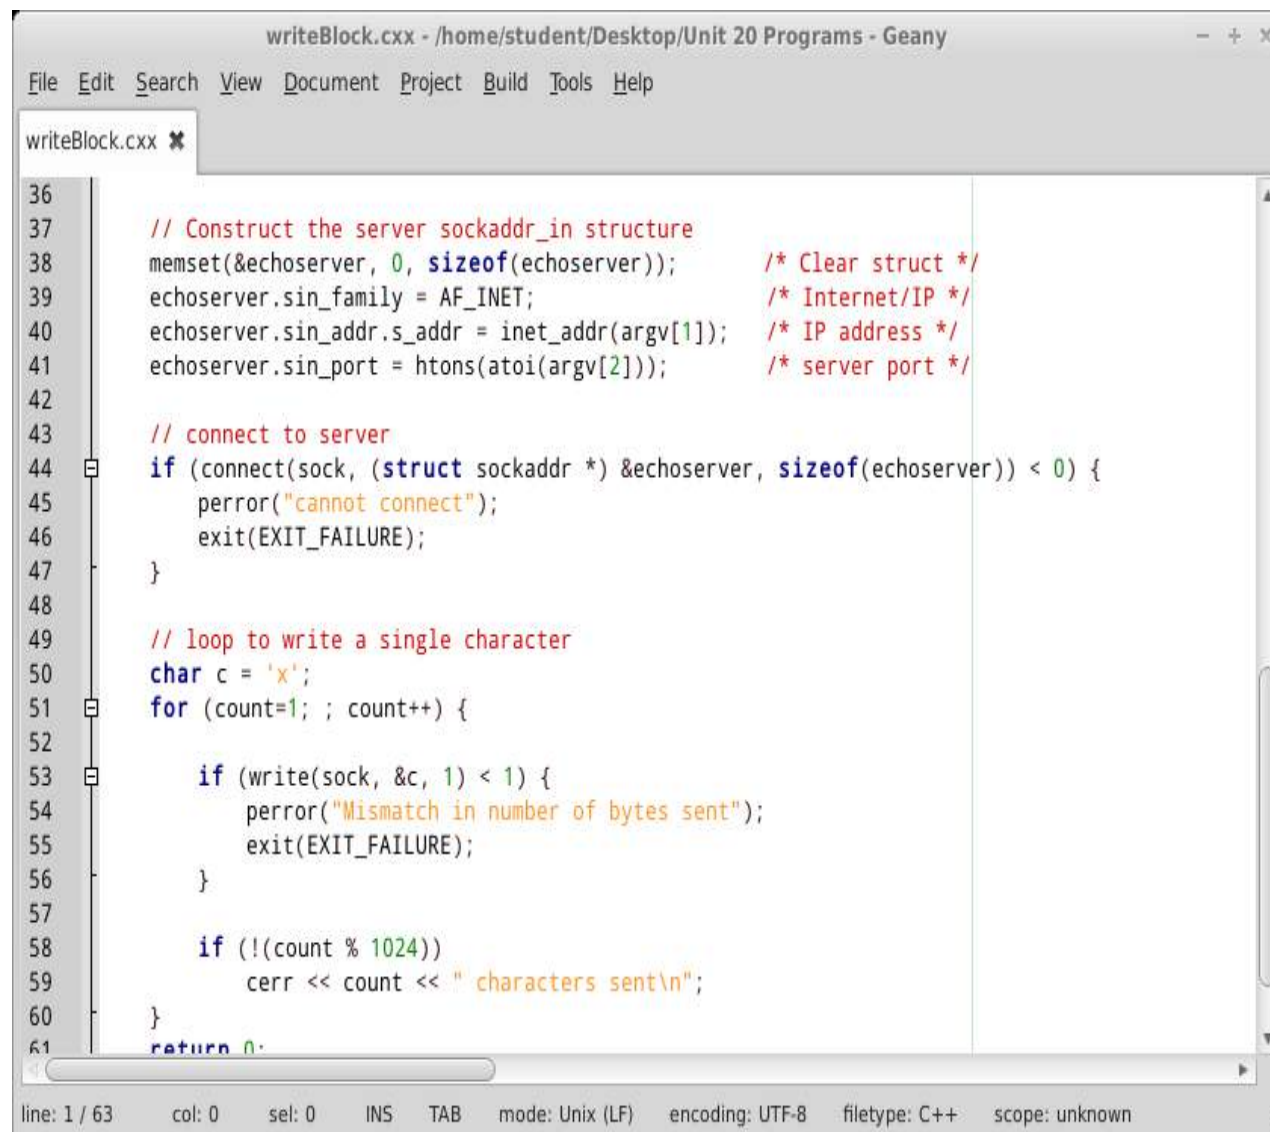


```
dummyTCPServer.cxx - /home/student/Desktop/Unit 20 Programs - Geany
File Edit Search View Document Project Build Tools Help

dummyTCPServer.cxx ✕

60
61 // Run until cancelled
62 while (true) {
63     connSock=accept(sock, NULL, NULL);
64     if (connSock < 0) {
65         perror("accept failed");
66         exit(EXIT_FAILURE);
67     }
68     // fork
69     if (fork()) { // parent process
70         close(connSock);
71     } else { // child process
72         cerr << "Waiting here: ";
73         cin >> buffer;
74
75         close(connSock);
76         exit(EXIT_SUCCESS);
77     }
78 }
79
80 close(sock);
81 return 0;
82
83
line: 73 / 83 col: 0 sel: 0 INS TAB mode: Unix (LF) encoding: UTF-8 filetype: C++ scope: main
```

Illustration: blocking mode



```
writeBlock.cxx - /home/student/Desktop/Unit 20 Programs - Geany
File Edit Search View Document Project Build Tools Help
writeBlock.cxx x
36
37 // Construct the server sockaddr_in structure
38 memset(&echoserver, 0, sizeof(echoserver)); /* Clear struct */
39 echoserver.sin_family = AF_INET; /* Internet/IP */
40 echoserver.sin_addr.s_addr = inet_addr(argv[1]); /* IP address */
41 echoserver.sin_port = htons(atoi(argv[2])); /* server port */
42
43 // connect to server
44 if (connect(sock, (struct sockaddr *) &echoserver, sizeof(echoserver)) < 0) {
45     perror("cannot connect");
46     exit(EXIT_FAILURE);
47 }
48
49 // loop to write a single character
50 char c = 'x';
51 for (count=1; ; count++) {
52
53     if (write(sock, &c, 1) < 1) {
54         perror("Mismatch in number of bytes sent");
55         exit(EXIT_FAILURE);
56     }
57
58     if (!(count % 1024))
59         cerr << count << " characters sent\n";
60 }
61 return 0;
```

line: 1 / 63 col: 0 sel: 0 INS TAB mode: Unix (LF) encoding: UTF-8 filetype: C++ scope: unknown

What to do when call blocks ?

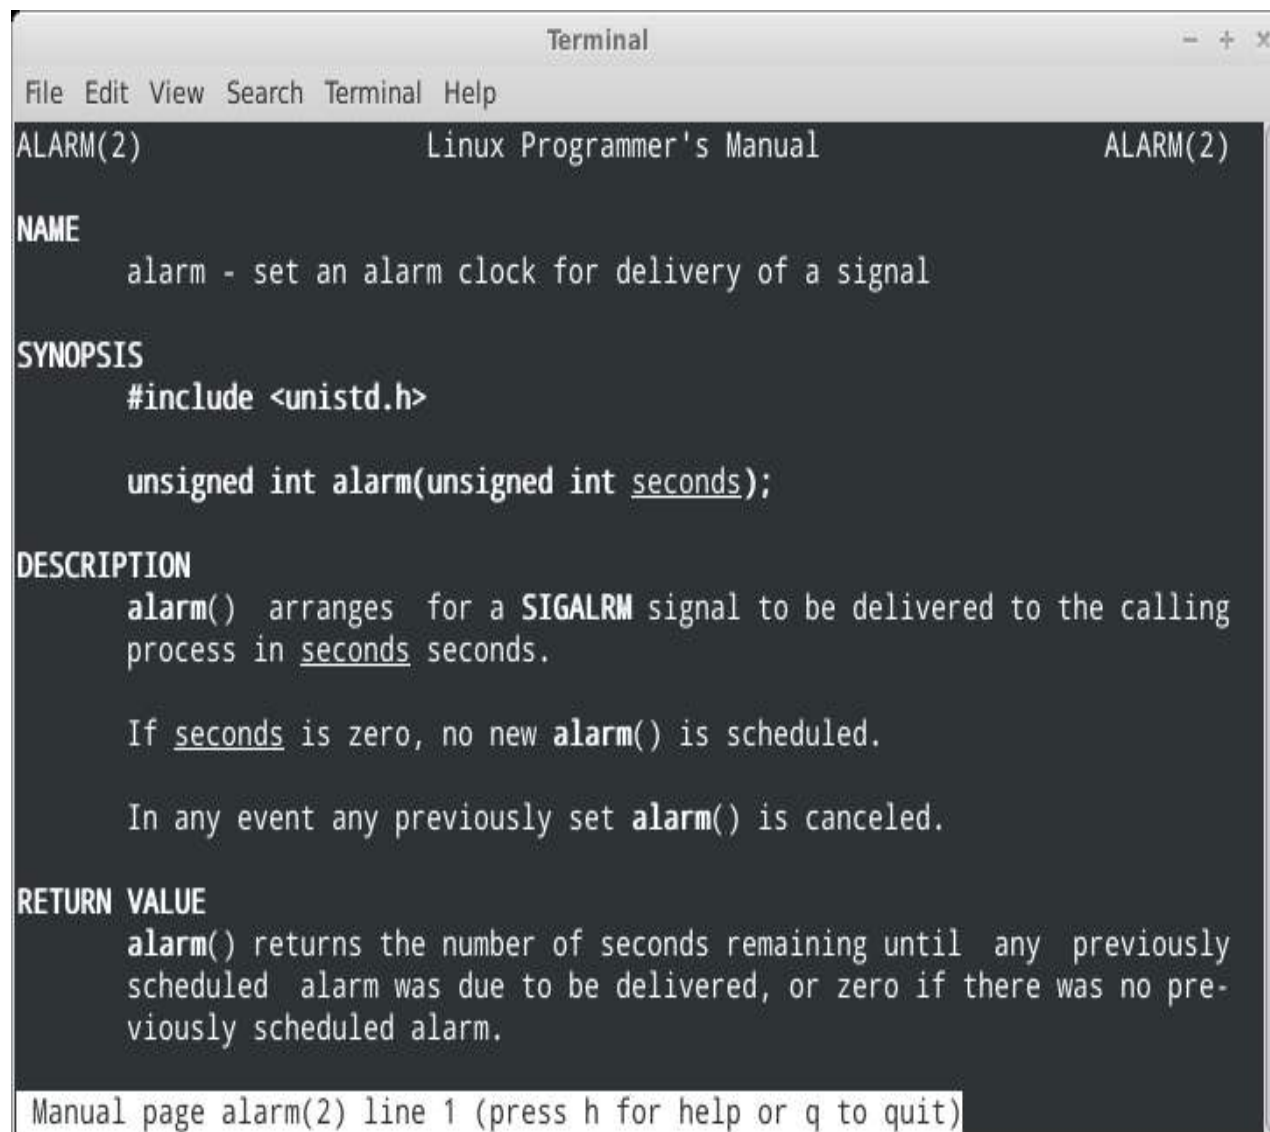
- Answer: wait

or: set alarm

– i.e. schedule signal to sent in future

- how:
 - set alarm timeout
 - program and install signal handler

System call: alarm



```
Terminal
File Edit View Search Terminal Help
ALARM(2) Linux Programmer's Manual ALARM(2)

NAME
    alarm - set an alarm clock for delivery of a signal

SYNOPSIS
    #include <unistd.h>

    unsigned int alarm(unsigned int seconds);

DESCRIPTION
    alarm() arranges for a SIGALRM signal to be delivered to the calling
    process in seconds seconds.

    If seconds is zero, no new alarm() is scheduled.

    In any event any previously set alarm() is canceled.

RETURN VALUE
    alarm() returns the number of seconds remaining until any previously
    scheduled alarm was due to be delivered, or zero if there was no pre-
    viously scheduled alarm.

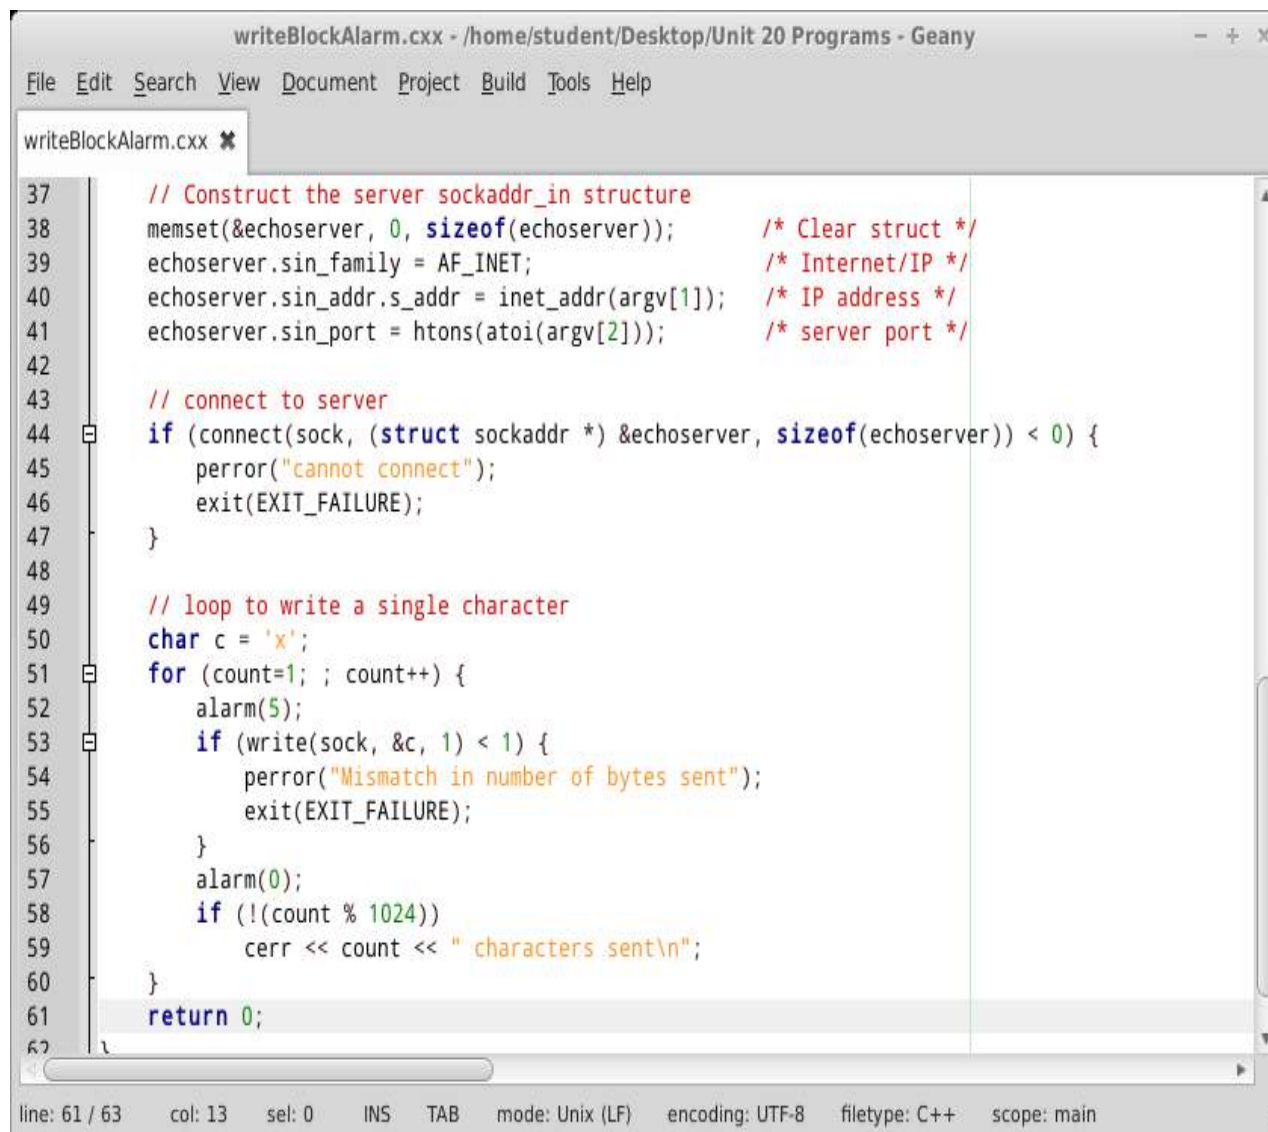
Manual page alarm(2) line 1 (press h for help or q to quit)
```

System call: alarm

`unsigned int alarm(unsigned int seconds)`

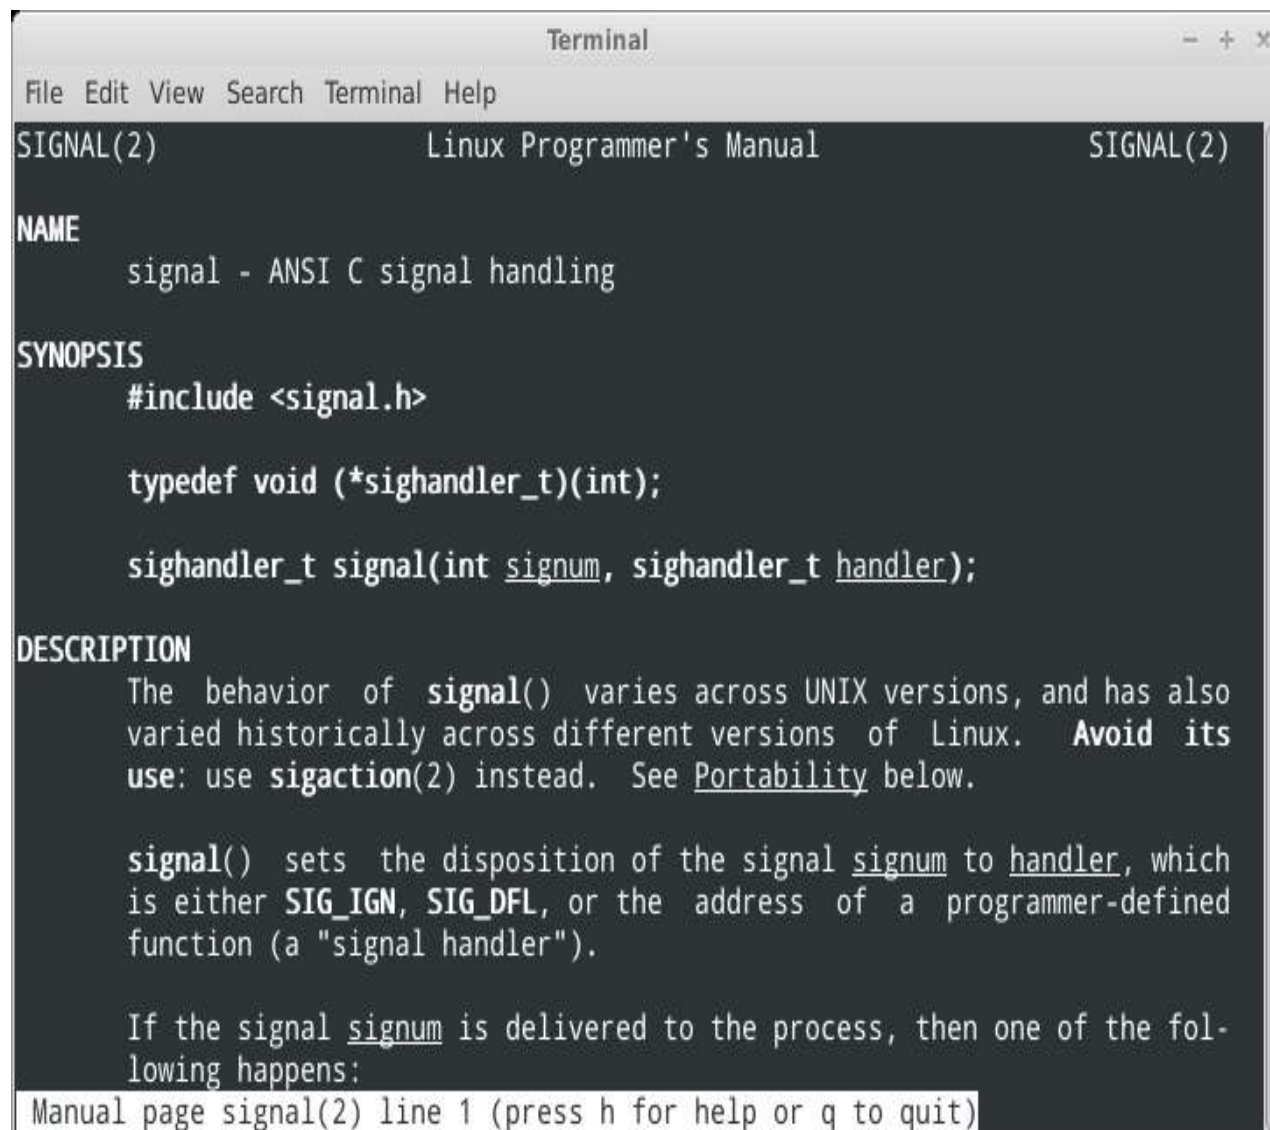
- arranges for a SIGALRM signal to be delivered to the calling process in `seconds` seconds
 - if `seconds` is zero, no new `alarm()` is scheduled
 - any previously set `alarm()` is canceled
- returns the number of seconds remaining until any previously scheduled alarm was due to be delivered
- zero if there was no previously scheduled alarm

Illustration: blocking mode



```
writeBlockAlarm.cxx - /home/student/Desktop/Unit 20 Programs - Geany
File Edit Search View Document Project Build Tools Help
writeBlockAlarm.cxx ✖
37 // Construct the server sockaddr_in structure
38 memset(&echoserver, 0, sizeof(echoserver)); /* Clear struct */
39 echoserver.sin_family = AF_INET; /* Internet/IP */
40 echoserver.sin_addr.s_addr = inet_addr(argv[1]); /* IP address */
41 echoserver.sin_port = htons(atoi(argv[2])); /* server port */
42
43 // connect to server
44 if (connect(sock, (struct sockaddr *) &echoserver, sizeof(echoserver)) < 0) {
45     perror("cannot connect");
46     exit(EXIT_FAILURE);
47 }
48
49 // loop to write a single character
50 char c = 'x';
51 for (count=1; ; count++) {
52     alarm(5);
53     if (write(sock, &c, 1) < 1) {
54         perror("Mismatch in number of bytes sent");
55         exit(EXIT_FAILURE);
56     }
57     alarm(0);
58     if (!(count % 1024))
59         cerr << count << " characters sent\n";
60 }
61 return 0;
62
line: 61 / 63 col: 13 sel: 0 INS TAB mode: Unix (LF) encoding: UTF-8 filetype: C++ scope: main
```

System call: signal



```
Terminal
File Edit View Search Terminal Help
SIGNAL(2)          Linux Programmer's Manual          SIGNAL(2)

NAME
    signal - ANSI C signal handling

SYNOPSIS
    #include <signal.h>

    typedef void (*sighandler_t)(int);

    sighandler_t signal(int signum, sighandler_t handler);

DESCRIPTION
    The behavior of signal() varies across UNIX versions, and has also
    varied historically across different versions of Linux. Avoid its
    use: use sigaction(2) instead. See Portability below.

    signal() sets the disposition of the signal signum to handler, which
    is either SIG_IGN, SIG_DFL, or the address of a programmer-defined
    function (a "signal handler").

    If the signal signum is delivered to the process, then one of the fol-
    lowing happens:

Manual page signal(2) line 1 (press h for help or q to quit)
```

System call: signal

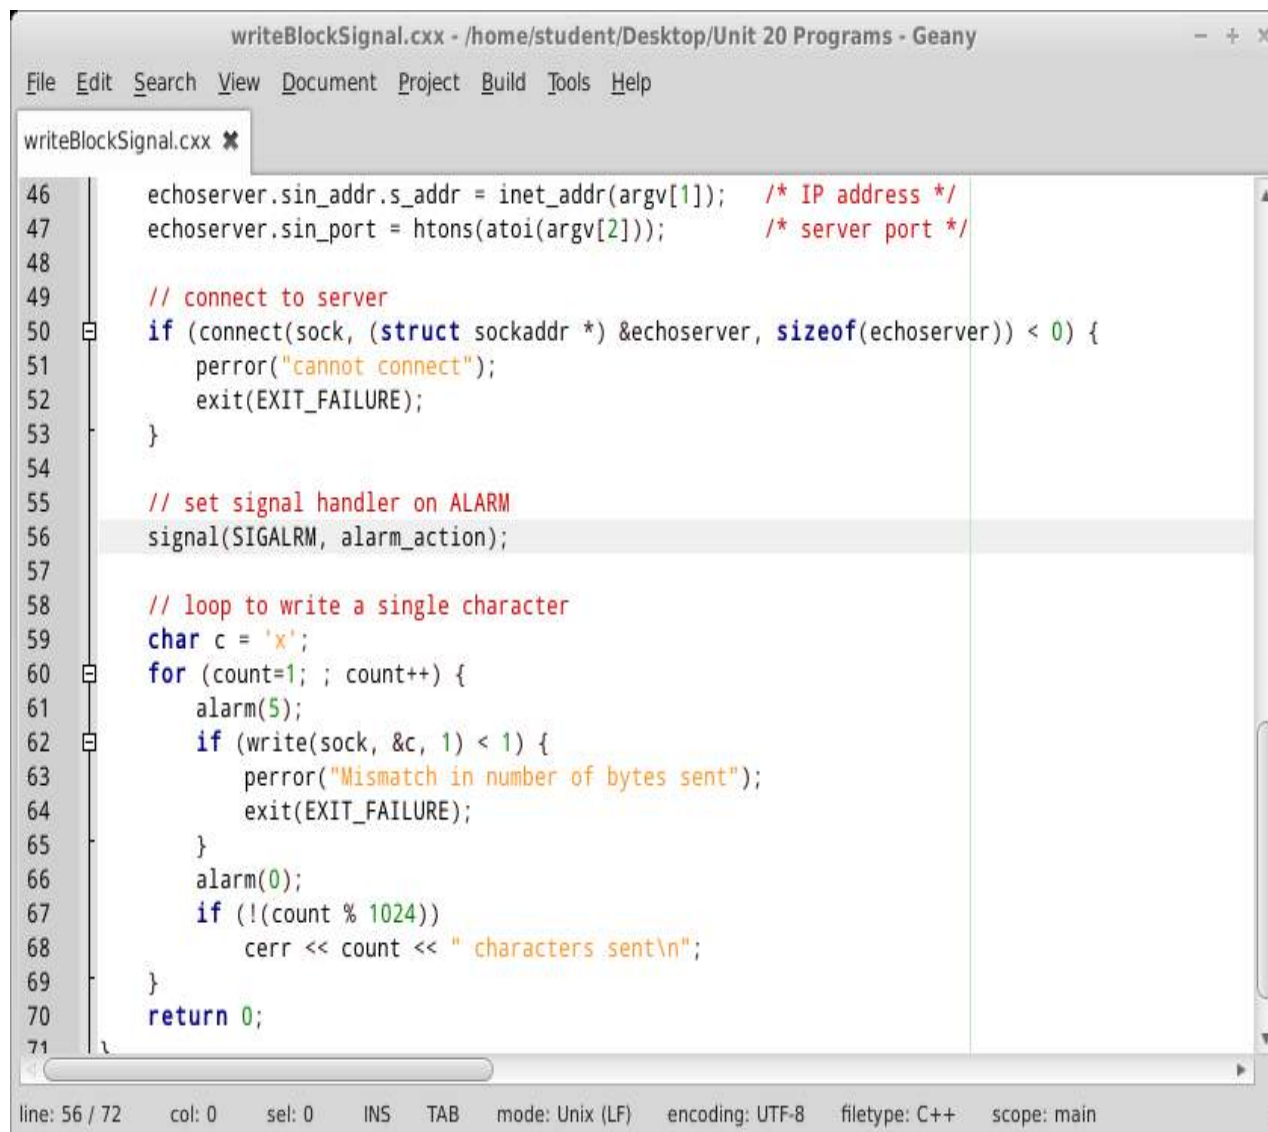
```
typedef void (*sighandler_t) (int);  
sighandler_t signal(int signum, sighandler_t handler)
```

- sets the disposition of the signal **signum** to **handler**
- **handler** can be:
 - **SIG_IGN** to ignore signal
 - **SIG_DFL** to restore default
 - address of a programmer-defined function
- returns the previous value of the signal handler

Signal handler: detail

```
void alarm_action(int s) {  
    cerr << "write blocked after " << count << " characters\n";  
    exit(EXIT_SUCCESS);  
}  
  
// set signal handler on ALARM  
signal(SIGALRM, alarm_action);
```

Illustration: blocking mode w/signal



```
writeBlockSignal.cxx - /home/student/Desktop/Unit 20 Programs - Geany
File Edit Search View Document Project Build Tools Help
writeBlockSignal.cxx *
46 echoserver.sin_addr.s_addr = inet_addr(argv[1]); /* IP address */
47 echoserver.sin_port = htons(atoi(argv[2])); /* server port */
48
49 // connect to server
50 if (connect(sock, (struct sockaddr *) &echoserver, sizeof(echoserver)) < 0) {
51     perror("cannot connect");
52     exit(EXIT_FAILURE);
53 }
54
55 // set signal handler on ALARM
56 signal(SIGALRM, alarm_action);
57
58 // loop to write a single character
59 char c = 'x';
60 for (count=1; ; count++) {
61     alarm(5);
62     if (write(sock, &c, 1) < 1) {
63         perror("Mismatch in number of bytes sent");
64         exit(EXIT_FAILURE);
65     }
66     alarm(0);
67     if (!(count % 1024))
68         cerr << count << " characters sent\n";
69 }
70 return 0;
71
line: 56 / 72 col: 0 sel: 0 INS TAB mode: Unix (LF) encoding: UTF-8 filetype: C++ scope: main
```


TCP socket: non-blocking mode

- calls returns immediately, producing full, partial or no result
- socket can be in blocking or non-blocking mode
- determined by flag associated with socket descriptor
- flag can be set via `fcntl` system call

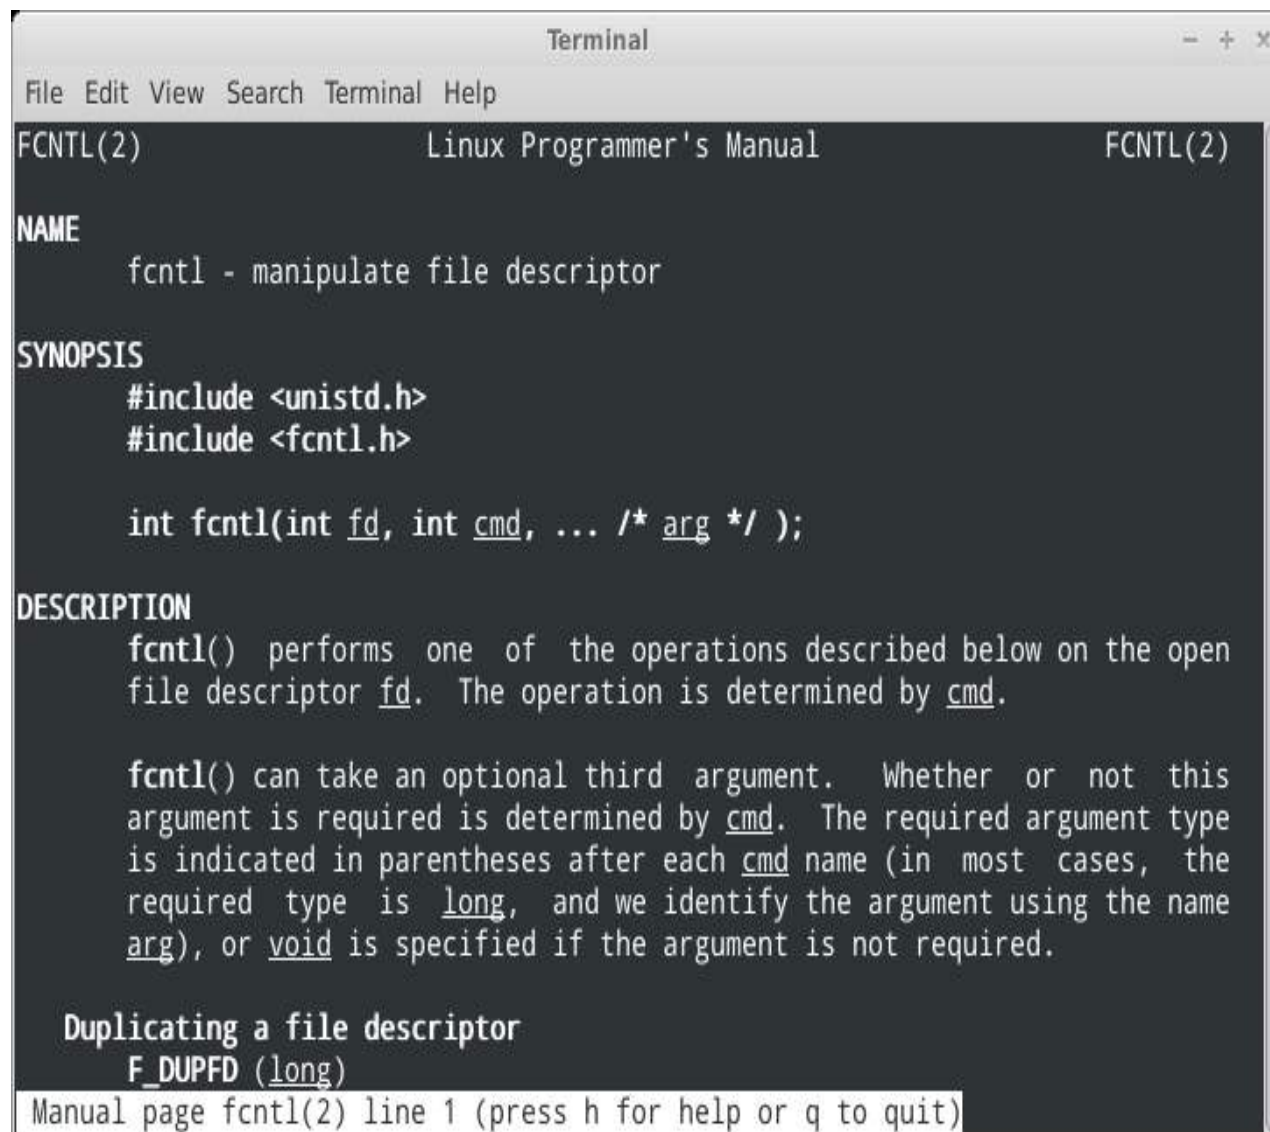
Non-Blocking mode: call behavior

- **accept**
 - when there are connections to accept, accept behaves as usual
 - if there are no pending connections, returns `-1` with `errno` set to `EWOULDBLOCK`
- **connect**
 - when connection to a listening server can be established, connect behaves as usual
 - If connection cannot be established, connect returns `-1` with `errno` set to `EINPROGRESS`

Non-Blocking mode: call behavior

- read
 - when there is data to read, read behaves as usual
 - if there is no data, read returns `-1` with `errno EWOULDBLOCK`
- write
 - when buffer space is available, write behaves as usual
 - if buffer becomes full, write returns number of bytes written
 - if buffer is full, write returns `-1` with `errno EWOULDBLOCK`

System call: fcntl



```
Terminal
File Edit View Search Terminal Help
FCNTL(2) Linux Programmer's Manual FCNTL(2)

NAME
fcntl - manipulate file descriptor

SYNOPSIS
#include <unistd.h>
#include <fcntl.h>

int fcntl(int fd, int cmd, ... /* arg */ );

DESCRIPTION
fcntl() performs one of the operations described below on the open
file descriptor fd. The operation is determined by cmd.

fcntl() can take an optional third argument. Whether or not this
argument is required is determined by cmd. The required argument type
is indicated in parentheses after each cmd name (in most cases, the
required type is long, and we identify the argument using the name
arg), or void is specified if the argument is not required.

Duplicating a file descriptor
F_DUPFD (long)

Manual page fcntl(2) line 1 (press h for help or q to quit)
```

System call: fcntl

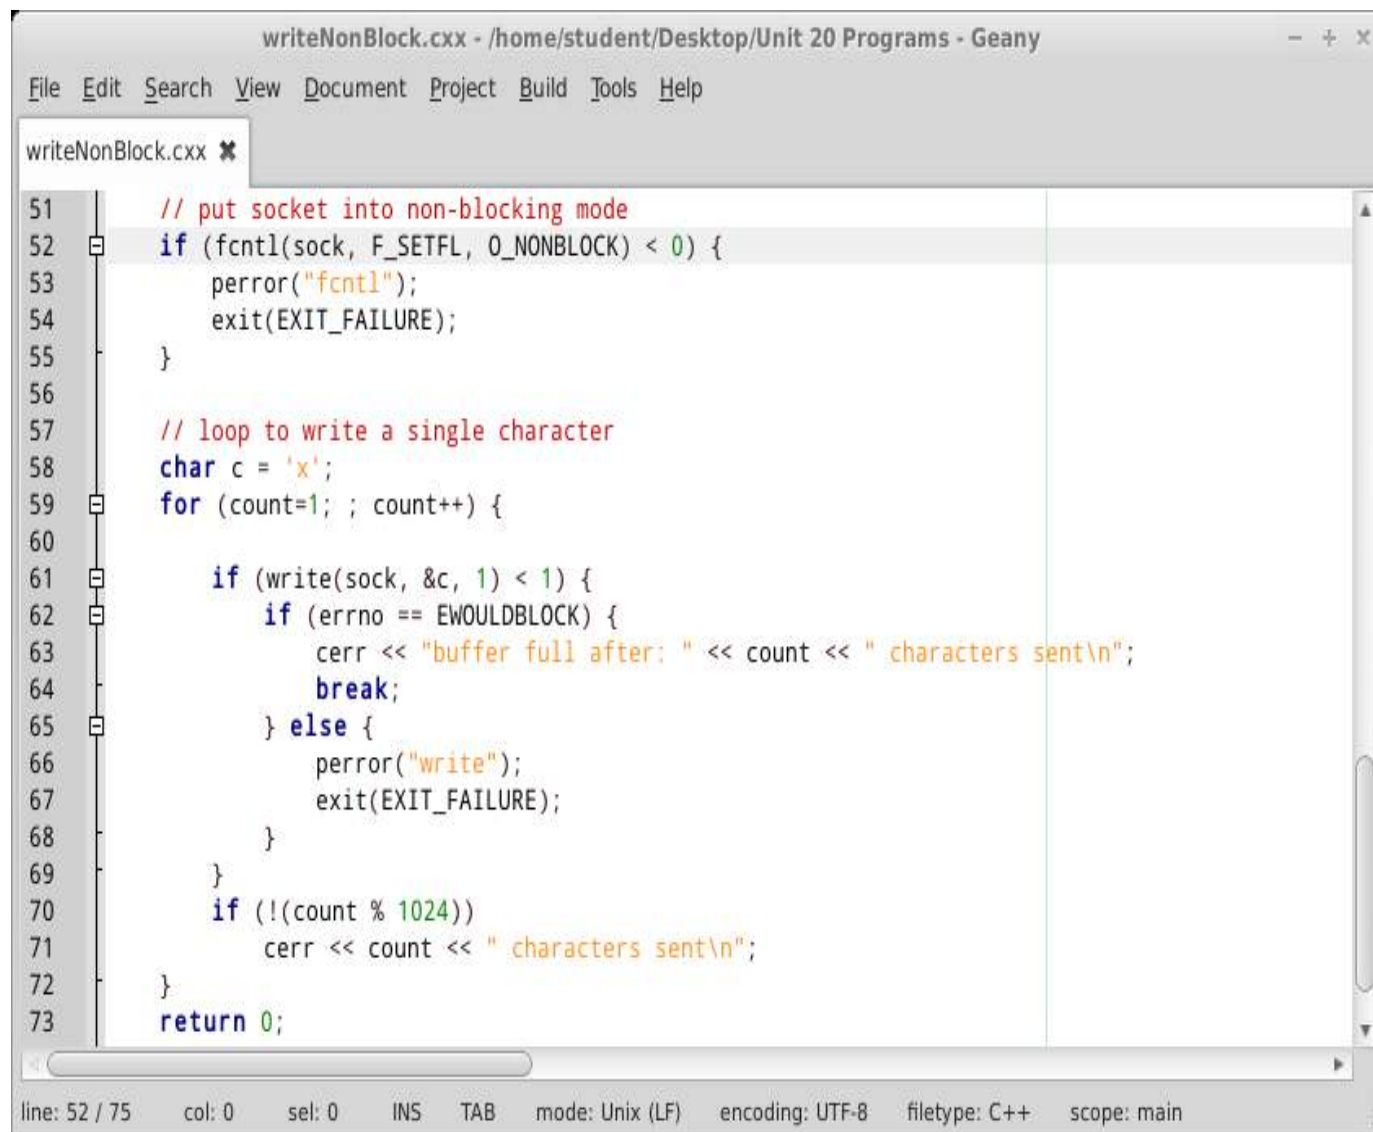
```
int fcntl(int fd, int cmd, ... /* arg */ )
```

- performs operation `cmd` on the open descriptor `fd`
- `cmd` to set non-blocking mode for socket:

`F_SETFL, O_NONBLOCK`

- returns -1 on error

Illustration: non-blocking write



The screenshot shows a Geany IDE window titled "writeNonBlock.cxx - /home/student/Desktop/Unit 20 Programs - Geany". The code is written in C++ and demonstrates how to set a socket to non-blocking mode and then attempt to write a single character. The code uses `fcntl` to set `O_NONBLOCK` on a socket. It then enters a loop to write a character 'x'. If the write operation returns less than 1, it checks for `EWOULDBLOCK`. If it is `EWOULDBLOCK`, it prints a message indicating the buffer is full after a certain number of characters and breaks the loop. Otherwise, it prints the number of characters sent. The code also includes error handling for `fcntl` and `write` failures.

```
51 // put socket into non-blocking mode
52 if (fcntl(sock, F_SETFL, O_NONBLOCK) < 0) {
53     perror("fcntl");
54     exit(EXIT_FAILURE);
55 }
56
57 // loop to write a single character
58 char c = 'x';
59 for (count=1; ; count++) {
60
61     if (write(sock, &c, 1) < 1) {
62         if (errno == EWOULDBLOCK) {
63             cerr << "buffer full after: " << count << " characters sent\n";
64             break;
65         } else {
66             perror("write");
67             exit(EXIT_FAILURE);
68         }
69     }
70     if (!(count % 1024))
71         cerr << count << " characters sent\n";
72 }
73 return 0;
```

line: 52 / 75 col: 0 sel: 0 INS TAB mode: Unix (LF) encoding: UTF-8 filetype: C++ scope: main

Summary

- TCP programming
- socket behavior
 - blocking vs. non-blocking
- signal handler